

In [24]:

```
# %% [markdown]
# # Neuromorphic KWS: PyTorch SNN + Loihi Emulator
#
# This notebook:
# 1. Loads the SpeechCommands sample data (local folder).
# 2. Uses the same MFCC preprocessing as the baseline CNN.
# 3. Rebuilds the CNN + SNN architecture in PyTorch.
# 4. Loads your trained weights from `saved_models/`.
# 5. Runs PyTorch SNN inference for a test example.
# 6. Builds a small Nengo+Loihi network for a toy demo (later cells).

# %%
import os
import numpy as np
import torch
import torch.nn as nn
import torch.nn.functional as F

from pathlib import Path

import torchaudio
import torchaudio.transforms as T

import nengo
import nengo_loihi

print("PyTorch :", torch.__version__)
print("Torchaudio :", torchaudio.__version__)
print("Nengo :", nengo.__version__)
print("Nengo Loihi :", nengo_loihi.__version__)

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Device:", device)
```

PyTorch : 2.9.1
Torchaudio : 2.9.1
Nengo : 4.1.0
Nengo Loihi : 1.1.0
Device: cpu

In [25]:

```
from pathlib import Path

# Project root = the folder ABOVE this notebook
NOTEBOOK_DIR = Path.cwd()
PROJECT_ROOT = NOTEBOOK_DIR.parent

print("Notebook is running from:", NOTEBOOK_DIR)
print("Project root resolved as:", PROJECT_ROOT)

DATA_DIR = PROJECT_ROOT / "sample_data" / "speech_commands_v0.02"
MODEL_DIR = PROJECT_ROOT / "saved_models"
```

```
print("DATA_DIR exists:", DATA_DIR.exists())
print("MODEL_DIR exists:", MODEL_DIR.exists())
```

Notebook is running from: /Users/maddy/Desktop/PLEP/Project/CS-576-Final-Project/loihi_emulator
Project root resolved as: /Users/maddy/Desktop/PLEP/Project/CS-576-Final-Project
DATA_DIR exists: True
MODEL_DIR exists: True

In [26]:

```
# %%
SAMPLE_RATE = 16000
N_MFCC = 40

import soundfile as sf
import librosa

def wav_to_mfcc(path: Path) -> torch.Tensor:
    """
    Load WAV using soundfile + compute MFCC using librosa.
    Avoids all TorchAudio/TorchCodec backends.
    """

    # ---- Load WAV using SoundFile ----
    waveform, sr = sf.read(str(path))          # numpy array
    if waveform.ndim > 1:                      # stereo -> mono
        waveform = waveform.mean(axis=1)

    waveform = waveform.astype("float32")

    # ---- Resample if needed ----
    if sr != SAMPLE_RATE:
        waveform = librosa.resample(waveform, orig_sr=sr, target_sr=SAMPLE_RATE)

    # ---- Compute MFCC using librosa ----
    mfcc = librosa.feature.mfcc(
        y=waveform,
        sr=SAMPLE_RATE,
        n_mfcc=N_MFCC,
        n_fft=400,
        hop_length=160,
        n_mels=40,
    ) # shape -> (40, T)

    # ---- Normalize ----
    mfcc = torch.tensor(mfcc, dtype=torch.float32)
    mfcc = (mfcc - mfcc.mean()) / (mfcc.std() + 1e-5)

    # ---- Optional clamp ----
    mfcc = torch.clamp(mfcc, -2.0, 2.0)

    return mfcc
```

```
In [31]: class CNN_KWS(nn.Module):
    def __init__(self, num_classes=6, flatten_dim=3840):
        super().__init__()
        self.flatten_dim = flatten_dim

        self.features = nn.Sequential(
            nn.Conv2d(1, 8, kernel_size=5, stride=1, padding=2),
            nn.ReLU(),
            nn.MaxPool2d(2),
            nn.Conv2d(8, 16, kernel_size=3, stride=1, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(2),
        )

        self.classifier = nn.Sequential(
            nn.Linear(self.flatten_dim, 64),
            nn.ReLU(),
            nn.Linear(64, num_classes),
        )

    def forward(self, x):
        x = x.unsqueeze(1) # [B, 1, 40, T]
        x = self.features(x)
        x = torch.flatten(x, 1)

        F = x.shape[1]
        if F > self.flatten_dim:
            x = x[:, :self.flatten_dim] # crop
        elif F < self.flatten_dim:
            pad = self.flatten_dim - F
            x = F.pad(x, (0, pad)) # pad

        return self.classifier(x)
```

```
In [32]: import snntorch as snn
from snntorch import surrogate

spike_grad = surrogate.fast_sigmoid()

class SNN_KWS(nn.Module):
    def __init__(self, base_cnn: CNN_KWS, num_steps: int = 50, beta: f
        super().__init__()
        self.num_steps = num_steps

        # Reuse trained CNN layers
        self.features = base_cnn.features
        self.fc1 = base_cnn.classifier[0] # Linear(flatten_dim ->
        self.fc2 = base_cnn.classifier[2] # Linear(64 -> num_class)

        # Get the expected flatten dimension from fc1
        self.flatten_dim = self.fc1.in_features # usually 3840 in yo
```

```
# SNN layers
self.lif1 = snn.Leaky(beta=beta, spike_grad=spike_grad)
self.lif2 = snn.Leaky(beta=beta, spike_grad=spike_grad)

def _fix_feature_dim(self, x):
    """
    Ensure x has shape [B, flatten_dim]
    Crop if too large, pad if too small.
    """
    B, F_in = x.shape

    if F_in > self.flatten_dim:
        return x[:, :self.flatten_dim]

    elif F_in < self.flatten_dim:
        pad = self.flatten_dim - F_in
        return F.pad(x, (0, pad))

    else:
        return x

def forward(self, x):
    """
    x: [B, 40, T]
    Output: [T, B, num_classes]
    """
    spk2_rec = []
    mem1 = self.lif1.init_leaky()
    mem2 = self.lif2.init_leaky()

    x = x.unsqueeze(1) # [B, 1, 40, T]

    for _ in range(self.num_steps):

        # CNN feature extractor
        cur = self.features(x)
        cur = torch.flatten(cur, 1) # [B, F_in]

        # Fix dimension mismatch
        cur = self._fix_feature_dim(cur) # -> [B, flatten_dim]

        # Fully connected -> LIF -> fc2 -> LIF
        cur = F.relu(self.fc1(cur))
        spk1, mem1 = self.lif1(cur, mem1)

        cur2 = self.fc2(spk1)
        spk2, mem2 = self.lif2(cur2, mem2)

        spk2_rec.append(spk2)

    return torch.stack(spk2_rec) # [T, B, C]
```

In [33]:

```
# %%
cnn_path = MODEL_DIR / "baseline_cnn_kws_vfinal.pt"
snn_path = MODEL_DIR / "snn_kws_model.pt"

print("CNN model file exists:", cnn_path.exists())
print("SNN model file exists:", snn_path.exists())

# Load checkpoint to read original flatten dim
cnn_state = torch.load(cnn_path, map_location=device)
flatten_dim_ckpt = cnn_state["classifier.0.weight"].shape[1]
print("Flatten dim in checkpoint =", flatten_dim_ckpt)

# Rebuild CNN using correct flatten dim
cnn_model = CNN_KWS(num_classes=6, flatten_dim=flatten_dim_ckpt).to(de
cnn_model.load_state_dict(cnn_state)
cnn_model.eval()
print("CNN model loaded successfully.")

# Build SNN model using CNN weights
snn_model = SNN_KWS(cnn_model, num_steps=50, beta=0.95).to(device)

# Load SNN weights if exist
if snn_path.exists():
    snn_state = torch.load(snn_path, map_location=device)
    snn_model.load_state_dict(snn_state)
    print("Loaded SNN weights from snn_kws_model.pt")
else:
    print("No separate SNN weights found; using CNN weights in SNN wra
snn_model.eval()
```

CNN model file exists: True
SNN model file exists: True
Flatten dim in checkpoint = 3840
CNN model loaded successfully.
Loaded SNN weights from snn_kws_model.pt

```
Out[33]: SNN_KWS(  
    (features): Sequential(  
        (0): Conv2d(1, 8, kernel_size=(5, 5), stride=(1, 1), padding=(2,  
        2))  
        (1): ReLU()  
        (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ce  
        il_mode=False)  
        (3): Conv2d(8, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1,  
        1))  
        (4): ReLU()  
        (5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ce  
        il_mode=False)  
    )  
    (fc1): Linear(in_features=3840, out_features=64, bias=True)  
    (fc2): Linear(in_features=64, out_features=6, bias=True)  
    (lif1): Leaky()  
    (lif2): Leaky()  
)
```

```
In [34]: # %%  
import random  
  
# restrict to the same keywords you used in training  
CLASSES = ["yes", "no", "go", "stop", "down", "up"]  
  
def pick_example(label: str = None) -> Path:  
    labels = [d for d in DATA_DIR.iterdir() if d.is_dir() and d.name in CLASSES]  
    if label is None:  
        label_dir = random.choice(labels)  
    else:  
        label_dir = DATA_DIR / label  
    files = sorted([p for p in label_dir.glob("*.wav")])  
    return random.choice(files)  
  
example_path = pick_example("yes") # or None for random  
print("Example file:", example_path)  
  
mfcc = wav_to_mfcc(example_path) # [40, T]  
mfcc_batch = mfcc.unsqueeze(0).to(device) # [1, 40, T]  
  
with torch.no_grad():  
    out_TBC = snn_model(mfcc_batch) # [T, 1, C]  
    logits = out_TBC.sum(dim=0).squeeze(0) # [C]  
    pred_idx = logits.argmax().item()  
  
print("Predicted class:", CLASSES[pred_idx])  
  
Example file: /Users/maddy/Desktop/PLEP/Project/CS-576-Final-Project/sa  
mple_data/speech_commands_v0.02/yes/2748cce7_nohash_0.wav  
Predicted class: up
```

```
In [35]: # %%
```

```

def extract_cnn_features(x: torch.Tensor, model: CNN_KWS) -> torch.Tensor
    ....
    x: [B, 40, T]
    Returns: [B, 64] feature vector before final linear layer.
    ....
    with torch.no_grad():
        x_ = x.unsqueeze(1) # [B, 1, 40, T]
        h = model.features(x_)
        h = torch.flatten(h, 1)

        # --- FIX FEATURE DIMENSION ---
        flatten_dim = model.classifier[0].in_features # e.g., 3840
        F_in = h.shape[1]

        if F_in > flatten_dim:
            h = h[:, :flatten_dim]
        elif F_in < flatten_dim:
            pad = flatten_dim - F_in
            h = F.pad(h, (0, pad))

        # first FC + ReLU
        fc1 = model.classifier[0]
        h = F.relu(fc1(h))

    return h # [B, 64]

# Get feature vector for the same example
feat = extract_cnn_features(mfcc_batch, cnn_model) # [1, 64]
feat_np = feat.cpu().numpy().flatten()

# Get fc2 weights and bias
fc2 = cnn_model.classifier[2]
W = fc2.weight.detach().cpu().numpy() # [6, 64]
b = fc2.bias.detach().cpu().numpy() # [6]

W_T = W # [64, 6]

num_classes = W.shape[0]
print("Feature dim:", feat_np.shape, "| num_classes:", num_classes)

```

Feature dim: (64,) | num_classes: 6

```

In [36]: # %%
with nengo.Network(seed=0) as net:

    # 1) Input feature vector (64-dim)
    inp = nengo.Node(feat_np) # constant for this example

    # 2) LIF Ensemble representing 64 features
    ens = nengo.Ensemble(
        n_neurons=64,
        dimensions=64,

```

```
    neuron_type=nengo.LIF(),      # Loihi-friendly neuron
    max_rates=nengo.dists.Uniform(100, 200),
)

# 3) Output node (6 class logits)
out = nengo.Node(size_in=num_classes)

# 4) Connect input → ensemble
nengo.Connection(inp, ens, synapse=None)

# 5) Connect ensemble.neurons → output using CNN weights
#     W_T is shape (6,64), perfect for (post=6, pre=64)
scale = 0.01
nengo.Connection(
    ens.neurons,
    out,
    transform=W_T * scale,
    synapse=0.005,
)

# 6) Probe the output
p_out = nengo.Probe(out, synapse=0.01)

# --- Run Loihi emulator ---
with nengo_loihi.Simulator(net) as sim:
    sim.run(0.1) # simulate 100 ms

# --- Extract results ---
logits_loihi = sim.data[p_out][-1]
pred_loihi = int(np.argmax(logits_loihi))

print("Logits (Loihi):", logits_loihi)
print("Predicted class (Loihi):", CLASSES[pred_loihi])
```

```
Logits (Loihi): [-0.44104002 -0.68619463 -0.00359438 -0.45080814 -0.642
58878 -0.19853477]
Predicted class (Loihi): go
```

In []:

```
# %%
import nengo
import nengo_loihi
import numpy as np

def run_loihi_for_feature(feat_vec: np.ndarray, W: np.ndarray, sim_time
    ....
        Run the 64-D feature vector through a tiny Loihi-emulated classifi

        feat_vec: shape [64]
        W:         shape [6, 64] (same as cnn_model.classifier[2].weight)
        sim_time: simulation time in seconds (e.g., 0.1 = 100 ms)
```

```
Returns:
    logits_loihi: numpy array [6]
    ....
    assert feat_vec.shape == (64,), f"Expected feature shape (64,), go
    assert W.shape[0] == 6, f"Expected W to have 6 rows (classes), got
    assert W.shape[1] == 64, f"Expected W to have 64 columns (features

    num_classes = W.shape[0]

    with nengo.Network(seed=0) as net:
        # 1) Constant input node: always outputs feat_vec
        inp = nengo.Node(output=lambda t: feat_vec)

        # 2) LIF ensemble representing the 64-D feature
        ens = nengo.Ensemble(
            n_neurons=64,
            dimensions=64,
            neuron_type=nengo.LIF(),
        )

        # 3) Output node: 6-D classification logits
        out = nengo.Node(size_in=num_classes)

        # 4) Input → ensemble (identity)
        nengo.Connection(inp, ens, synapse=None)

        # 5) Ensemble neurons → output using W (6 x 64)
        nengo.Connection(
            ens.neurons,
            out,
            transform=W,
            synapse=0.01,
        )

        # 6) Probe output with a small synapse for smoothing
        p_out = nengo.Probe(out, synapse=0.01)

    # Run Loihi emulator
    with nengo_loihi.Simulator(net) as sim:
        sim.run(sim_time)
        logits_loihi = sim.data[p_out][-1] # last timestep

    # Clean up any numerical weirdness
    logits_loihi = np.nan_to_num(logits_loihi)

    return logits_loihi
```

```
In [38]: # %%
from typing import Tuple

def eval_loihi_classifier(
    loader,
```

```
cnn_model: CNN_KWS,
W: np.ndarray,
device: torch.device,
max_samples: int = 50,
sim_time: float = 0.1,
) -> Tuple[float, float, int]:
    """
    Compare CNN classifier vs Loihi emulator on a subset of the test set.

    Returns:
        cnn_acc: CNN head accuracy on the subset
        loihi_acc: Loihi classifier accuracy on the subset
        total: number of samples evaluated
    """
    cnn_model.eval()

    total = 0
    correct_cnn = 0
    correct_loihi = 0

    for mfcc_batch, y_batch in loader:
        mfcc_batch = mfcc_batch.to(device)
        y_batch_np = y_batch.numpy()

        with torch.no_grad():
            # [B,64] features from CNN
            feats = extract_cnn_features(mfcc_batch, cnn_model)
            # CNN head logits [B,6]
            fc2 = cnn_model.classifier[2]
            logits_cnn = fc2(feats)
            preds_cnn = logits_cnn.argmax(dim=1).cpu().numpy()

            batch_size = feats.size(0)

            for i in range(batch_size):
                feat_np = feats[i].cpu().numpy()
                label = int(y_batch_np[i])

                # PyTorch CNN prediction
                if preds_cnn[i] == label:
                    correct_cnn += 1

                # Loihi prediction
                logits_loihi = run_loihi_for_feature(
                    feat_vec=feat_np,
                    W=W,
                    sim_time=sim_time,
                )
                pred_loihi = int(np.argmax(logits_loihi))

                if pred_loihi == label:
                    correct_loihi += 1
```

```
    total += 1
    if total >= max_samples:
        cnn_acc = correct_cnn / total
        loihi_acc = correct_loihi / total
        return cnn_acc, loihi_acc, total

    cnn_acc = correct_cnn / max(total, 1)
    loihi_acc = correct_loihi / max(total, 1)
    return cnn_acc, loihi_acc, total
```

In [39]:

```
# %%
import torchaudio
import torch.nn.functional as F
from torch.utils.data import DataLoader, Dataset
from pathlib import Path
import soundfile as sf

DATA_DIR = PROJECT_ROOT / "sample_data/speech_commands_v0.02"

CLASSES = ["yes", "no", "go", "stop", "down", "up"]
SAMPLE_RATE = 16000
N_MFCC = 40

mfcc_transform = torchaudio.transforms.MFCC(
    sample_rate=SAMPLE_RATE,
    n_mfcc=N_MFCC,
    melkwargs={
        "n_fft": 400,
        "hop_length": 160,
        "n_mels": 40,
        "center": False,
    },
)

# -----
# FIXED: safe MFCC loader (soundfile)
# -----
def wav_to_mfcc(path: Path) -> torch.Tensor:
    waveform, sr = sf.read(str(path))
    waveform = torch.tensor(waveform).float().unsqueeze(0)

    if sr != SAMPLE_RATE:
        waveform = torchaudio.functional.resample(waveform, sr, SAMPLE_RATE)

    mfcc = mfcc_transform(waveform).squeeze(0)
    mfcc = (mfcc - mfcc.mean()) / (mfcc.std() + 1e-5)
    return mfcc

# -----
# FIXED: Dataset now receives file LIST, not folder
```

```
# -----
# class KWS_Dataset(Dataset):
#     def __init__(self, file_list, classes):
#         self.files = file_list
#         self.classes = classes

#     def __len__(self):
#         return len(self.files)

#     def __getitem__(self, idx):
#         path = self.files[idx]
#         mfcc = wav_to_mfcc(path)
#         label = path.parent.name
#         y = self.classes.index(label)
#         return mfcc, y

# -----
# FIX: Collect all wav files
# -----
file_list = []
for label in CLASSES:
    folder = DATA_DIR / label
    file_list.extend(sorted(folder.glob("*.wav")))

print("Total WAV files loaded:", len(file_list))

# -----
# Create test_loader properly
# -----
test_dataset = KWS_Dataset(file_list, CLASSES)

def pad_collate(batch):
    xs, ys = zip(*batch)
    max_t = max(x.shape[1] for x in xs)
    xs = [F.pad(x, (0, max_t - x.shape[1])) for x in xs]
    xs = torch.stack(xs)
    ys = torch.tensor(ys)
    return xs, ys

test_loader = DataLoader(
    test_dataset,
    batch_size=1,
    shuffle=True,
    collate_fn=pad_collate,
)

print("Loaded test samples:", len(test_dataset))
```

Total WAV files loaded: 23377
Loaded test samples: 23377

```
In [40]: # %%
max_samples = 50    # adjust if you want more / fewer
sim_time = 0.1      # 100 ms per sample

cnn_acc, loihi_acc, total = eval_loihi_classifier(
    loader=test_loader,
    cnn_model=cnn_model,
    W=W,
    device=device,
    max_samples=max_samples,
    sim_time=sim_time,
)

print(f"Evaluated on {total} test samples")
print(f"CNN head accuracy: {cnn_acc*100:.2f}%")
print(f"Loihi classifier acc: {loihi_acc*100:.2f}%")
```

```
Evaluated on 50 test samples
CNN head accuracy: 10.00%
Loihi classifier acc: 12.00%
```

```
In [ ]:
```